

# Saliency-Driven Gaze Simulation with Joint Embedding Predictive Architectures

**Author:** Arash Khosropour

**Supervisors:** Prof. Giuseppe Boccignone

**Institution:** Università degli Studi di Milano

---

## 1. Introduction

Understanding where humans focus their attention within dynamic visual scenes is a foundational milestone in natural interaction research. Human gaze patterns, comprising sequences of precise fixations and rapid shifts known as saccades, reveal how the brain prioritizes, filters, and processes complex sensory information. Modeling these dynamics effectively is crucial for computational cognitive science.

This report presents a thorough structural evaluation and reverse engineering of SaccadeJEPA, a published architecture designed around predictive vision principles. While the architecture claims to simulate realistic saccadic eye movements natively, our investigation reveals a critical structural limitation: as originally implemented, the model selects gaze targets uniformly at random, serving purely as a data-augmentation tool rather than an active gaze model. To address this, we introduce and evaluate a principled content-aware alternative. By extracting latent features from frozen self-supervised and supervised backbones and training a lightweight prediction head on real human eye-tracking data, we demonstrate that integrating true visual saliency transforms the system into a high-fidelity gaze simulator.

## 2. Joint Embedding Predictive Architectures (JEPA)

The Joint Embedding Predictive Architecture (JEPA), pioneered by LeCun (2022), introduces a paradigm shift in self-supervised representation learning. Traditional generative architectures, such as Masked Autoencoders (MAE), train networks to reconstruct omitted data directly in the high-dimensional pixel space. This pixel-level reconstruction often forces the model to allocate significant capacity to modeling low-level, high-frequency noise and pixel-level statistics that lack semantic value.

In contrast, a JEPA operates entirely within a latent embedding space. The architecture avoids pixel-level reconstruction by training a predictor network to predict the embedding of a target patch from the embedding of a context patch. I-JEPA (Assran et al., 2023) adapts this concept to the vision domain using a Vision Transformer (ViT) backbone. By masking vast semantic regions of an input image, encoding only the visible context patches, and predicting the representations of the missing regions, the model learns highly robust, abstract semantic features. Crucially, the target encoder is updated as an Exponential Moving Average (EMA) of the context encoder, providing stable, non-collapsing targets without requiring negative samples or contrastive pairs.

### 3. SaccadeJEPA: Reverse Engineering and Findings

SaccadeJEPA (LumenPallidium, 2023) attempts to extend the core JEPA framework to the domain of human active vision and saccadic simulation. Rather than applying traditional rectangular patch masks, it creates a dual-view paradigm from a single input image. The network generates a shifted view by applying a random affine transformation (combining spatial translation and rotation), and sets a central crop of the source image as the target reference view. The network's core task is to predict how the latent representation shifts as a direct function of that affine spatial transformation.

A deep review of the codebase reveals that the mechanism responsible for selecting gaze changes (located in `saccade.py` line 66) is implemented as follows:

```
rotations      = torch.empty(b, 1).uniform_(-max_rotation_rad, max_rotation_rad)
translations   = torch.empty(b, 2).uniform_(-max_translation_frac,
max_translation_frac)
```

#### Critical System Finding

This implementation reveals that the "attention mechanism" within SaccadeJEPA is entirely governed by a uniform pseudo-random number generator. There is no content-driven selection, face tracking, or semantic weighting. The model does not look at areas because they are interesting; instead, the spatial translation represents a self-supervised data-augmentation technique (spatial jittering) to assist downstream feature representations. Consequently, it cannot natively simulate authentic gaze dynamics as required by biological visual systems.

### 4. Methodology: Saliency Prediction on the FIND Dataset

#### 4.1 The FIND Dataset

To establish a true, content-driven baseline, we utilize the FIND dataset provided by Prof. Alessandro D'Amelio. This dataset captures authentic human gaze interactions, containing eye-tracking recordings from 39 human observers watching 65 multi-face social videos sourced from YouTube and Youku at a resolution of  $1280 \times 720$  pixels. To evaluate models accurately without temporal data leakage, videos are split via a seeded random permutation (`seed=42`) into three distinct subsets: a training split (53 videos), a validation split (6 videos), and a testing split (6 videos). Splitting by video ID ensures the model never encounters adjacent frames of a test sequence during training.

#### 4.2 Ground Truth Heatmaps

Raw fixation coordinate points collected from observers are mapped onto a spatial probability distribution for each video frame. At each valid coordinate, a two-dimensional spatial Gaussian filter is applied with a scale factor of  $\sigma = 20$  pixels. The summation of these Gaussians across all active observers is normalized to sum to 1, yielding a continuous ground truth probability distribution matrix of  $224 \times 224$  pixels. This heatmap represents a precise empirical map of human visual focus.

#### 4.3 Saliency Sources Compared

Five distinct saliency sources were implemented, benchmarked, and evaluated:

- **Random Saliency (SaccadeJEPA Baseline):** A completely uniform map normalized to sum to 1. This directly models the random translation behavior native to SaccadeJEPA.

- **Center Bias (Parametric Baseline):** A fixed, non-learned 2D Gaussian distribution anchored directly at the center of the frame. This represents camera operator framing tendencies and standard viewing habits without any active image processing.
- **Spectral Residual (Classical Baseline):** A non-learned, signal-processing method based on local contrast and structural deviation from the average image statistics, implemented via OpenCV's StaticSaliencySpectralResidual (Hou & Zhang, 2007).
- **ResNetSaliency (Learned Supervised):** A supervised feature extractor utilizing a frozen ResNet-18 backbone pretrained on ImageNet, paired with a custom trained 2-layer convolutional head.
- **I-JEPA Saliency (Learned Self-Supervised):** A self-supervised feature extractor utilizing a frozen Vision Transformer (ViT-H/14) backbone from facebook/ijepa\_vith14\_1k, paired with an identical 2-layer convolutional prediction head.

#### 4.4 Network Architectures & Implementation Details

Both the ResNet-18 and I-JEPA models share an identical training strategy: the large backbone network is completely frozen to retain general features, and a small, randomly initialized convolutional head is optimized using the Kullback-Leibler (KL) divergence loss against human heatmaps over 20 epochs using the AdamW optimizer ( $\text{lr} = 1\text{e-}3$ ,  $\text{weight\_decay} = 1\text{e-}4$ ) and a Cosine Annealing learning rate schedule.

For ResNetSaliency, features are extracted at layer3, providing a  $14 \times 14$  spatial grid with 256 feature dimensions per cell. This resolution preserves optimal mid-level structural features (such as faces, hands, and object borders) while avoiding the excessive loss of spatial resolution found in layer4 ( $7 \times 7$ ). The head contains a  $3 \times 3$  convolution ( $256 \rightarrow 64$  channels), a GELU activation, and a  $1 \times 1$  convolution scaling down to 1 channel. This yields 147,585 trainable parameters.

For I-JEPA Saliency, the image is parsed into a sequence of 256 tokens via a  $14 \times 14$  pixel patch extraction step. Each patch yields a dense 1280-dimensional latent feature vector. These tokens are re-arranged in raster-scan order and reshaped into a  $16 \times 16 \times 1280$  spatial tensor. The head mirrors the ResNet architecture: a  $3 \times 3$  convolution ( $1280 \rightarrow 64$  channels), a GELU activation, and a  $1 \times 1$  convolution. This configuration yields 737,409 trainable parameters.

To ensure numerical stability and prevent noisy gradients stemming from sparse target structures, the KL-divergence loss function is evaluated directly at each model's native resolution ( $14 \times 14$  for ResNet,  $16 \times 16$  for I-JEPA) rather than at the full upsampled  $224 \times 224$  scale.

### 5. Quantitative Results and Benchmarking

The models were evaluated comprehensively against the held-out FIND test split (videos 070–075, encompassing 310 distinct validation frames) using three standard eye-tracking metrics: Area Under the ROC Curve (AUC-Judd version), Normalized Scanpath Saliency (NSS), and Pearson Correlation Coefficient (CC).

| Saliency Source               | Type                        | AUC $\uparrow$ | NSS $\uparrow$ | CC $\uparrow$ |
|-------------------------------|-----------------------------|----------------|----------------|---------------|
| Random (SaccadeJEPA baseline) | Stochastic Baseline         | 0.502          | 0.007          | 0.001         |
| Itti-Koch (Spectral Residual) | Classical Signal Processing | 0.688          | 0.394          | 0.136         |

|                 |                                        |              |              |              |
|-----------------|----------------------------------------|--------------|--------------|--------------|
| Center Bias     | Parametric Static Prior                | 0.772        | 1.045        | 0.458        |
| ResNetSaliency  | Supervised Frozen Backbone             | 0.881        | 2.525        | 0.723        |
| I-JEPA Saliency | <b>Self-Supervised Frozen Backbone</b> | <b>0.892</b> | <b>2.768</b> | <b>0.772</b> |

## 6. Discussion & Architectural Analysis

- **The Failure of Random Sampling:** The baseline SaccadeJEPA model scores identically to chance (AUC = 0.502, NSS  $\approx$  0), mathematically proving that it lacks intrinsic predictive power over human eye movements without modification.
- **Why Center Bias Outperforms Classical Saliency:** In the FIND dataset, Center Bias (AUC = 0.772) substantially outperforms the classical approach (AUC = 0.688). This occurs due to production bias cinematographers naturally frame subjects centrally, and human observers strongly anticipate this configuration. The classical detector isolates local contrast and high-frequency edges but misses these global composition rules.
- **The Self-Supervised Advantage of I-JEPA:** I-JEPA achieves the highest scores across all metrics (AUC = 0.892, NSS = 2.768, CC = 0.772). While ResNet-18 processes local features sequentially through convolutional filters, I-JEPA's self-attention mechanism enables long-range spatial context modeling. Every token embedding inherently encodes global spatial relationships, providing richer features for predicting human gaze.
- **The Role of the Head and Dataset Characteristics:** The performance gap between supervised ResNet and self-supervised I-JEPA is clear but relatively narrow. Because the FIND dataset is composed of social videos heavily dominated by faces and hands, the saliency task reduces primarily to a face localization objective. Since both ImageNet classification and self-supervised patch prediction develop high sensitivity to human features, both backbones solve this core task effectively.

## 7. Correcting SaccadeJEPA: Gaze-Loop Integration

To transform SaccadeJEPA into an authentic, content-driven gaze simulator, the uniform random translation mechanism can be replaced with a learned, saliency-guided sampling step. Instead of sampling translations blindly, the system treats the generated saliency map as a 2D multinomial probability distribution over the image space.

The complete, algorithmic integration process is defined as follows:

```
# 1. Generate content-aware saliency maps from input video frames
saliency_maps = ijepa_saliency_model(image_batch) # Shape: (B, 224, 224)

# 2. Flatten spatial dimensions into a 1D distribution per batch item
flat_distributions = saliency_maps.view(b, -1)

# 3. Perform a weighted multinomial draw to select a precise fixation cell
sampled_indices = torch.multinomial(flat_distributions,
num_samples=1).squeeze(1)
```

```
# 4. Map the 1D flat cell indices back to normalized device coordinates [-1, 1]
ys = (sampled_indices // 224).float() / 224.0 * 2.0 - 1.0
xs = (sampled_indices % 224).float() / 224.0 * 2.0 - 1.0

# 5. Overwrite the original random tensor with content-aware coordinates
translations = torch.stack([xs, ys], dim=1) # Shape: (B, 2)
```

By substituting this single code segment, the downstream JEPA framework—including the affine warping functions, positional embeddings, and predictive networks—remains intact, but is now driven by a content-aware prior. This integration reduces the architectural gap to perfect human gaze modeling by over 78% in terms of area under the curve.

## 8. Limitations & Future Outlook

While the proposed integration yields substantial performance improvements, several limitations remain. First, both models operate on isolated, single video frames, ignoring temporal dynamics and motion cues like optical flow, which heavily influence human gaze in dynamic environments. Second, the shallow, 2-layer convolutional head restricts complex spatial reasoning; replacing it with a multi-scale U-Net decoder with lateral skip connections would enhance spatial precision. Finally, the face-centric nature of the FIND dataset limits broader generalization. Fine-tuning the full backbone networks end-to-end on a more diverse, multi-domain eye-tracking dataset would be required to build a truly generalizable computational model of active vision.

---

## References

- Assran, M., Caron, M., Misra, I., Bojanowski, P., Joulin, A., Joulin, N., & LeCun, Y. (2023). Self-supervised learning from images with a joint-embedding predictive architecture. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 15619–15629.
- Bylinskii, Z., Judd, T., Oliva, A., Torralba, A., & Durand, F. (2019). What do different evaluation metrics tell us about saliency models? *IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)*, 41(3), 740–757.
- Hou, X., & Zhang, L. (2007). Saliency detection: A spectral residual approach. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 1–8.
- Itti, L., Koch, C., & Niebur, E. (1998). A model of saliency-based visual attention for rapid scene analysis. *IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)*, 20(11), 1254–1259.
- LeCun, Y. (2022). A path towards autonomous machine intelligence. *OpenReview Position Paper*.